



AI Builder Candidate Assignment Brief

Assessment exercise for KPMG AI Builder Senior Consultant and Manager candidates

Timebox: Maximum 3 hours total

Tools: Candidate's own tools, accounts, and AI tokens

Assessment-only assurance

This exercise is intended to evaluate your approach, judgment, and builder mindset. It is not a request to produce a formal hiring asset for immediate or current-cycle use. Nothing submitted as an assignment will be used by KPMG for anything other than the candidate's assessment.

Assignment

KPMG is hiring for an AI Builder role. The role is intentionally broad: strong candidates may come from software engineering, UX, MLOps, consulting, research, operations, product, or self-taught paths. What matters is the ability to enter an ambiguous enterprise context, identify where AI can create practical value, build something tangible, and explain tradeoffs clearly.

The Task

Build something you think would evaluate an AI Builder well. We are giving you a starting point, not a fixed problem: help the hiring team identify strong AI Builder candidates, using the AI Builder job description as your source material. Part of the task is deciding what “evaluate well” actually means, then building to your own definition. You choose the direction and the solution format. If you reframe the problem, tell us why.

Why the format is open

AI Builders join us from engineering, UX, MLOps, consulting, research, operations, product, and self-taught paths. A single fixed task would advantage one background and penalize the rest, and would tell us more about prior training than about how a candidate thinks. Leaving the format open lets you lean into your strengths and show your strongest work in your own domain. We want to see what you are actually good at, not watch you perform a skill you do not have. What every strong background shares is clear reasoning under ambiguity and constraint, and that is what we are looking for across all of them.

What we're really evaluating

We are not hiring you to build hiring tools. This exercise uses “evaluate an AI Builder well” as a deliberately open canvas, because the AI Builder role itself is open. You will routinely walk into ambiguous problems with no defined solution and decide what is worth building, and often what the real problem even is. The topic is the vehicle. What we assess is how



you get from a vague prompt to a concrete, defensible build: how you scope the problem, what you choose to optimize for, the choices you make, and the ones you deliberately do not.

What you can build

- A lightweight product, app, workflow, agent, or automation.
- An evaluation process, scoring or interview tool, or simulation.
- A candidate workbench, reviewer workflow, prompt/evaluation framework, or prototype review process.
- Another practical artifact that helps evaluate AI Builder candidates.

There is no single correct answer. Choose the approach that best shows how you think and build.

Required submission

Submission item	Expectation
Artifact	A working prototype, workflow, process, tool, demo, rubric, simulation, or clearly described system that evaluates AI Builder candidates.
3-minute video	Use the first 1 minute to show what you built. Use the remaining 2 minutes to explain your thinking process, how you framed the problem, key choices, trade offs, risks, AI use, and what you would improve with more time.

Rules and guardrails

- Spend no more than 3 hours total, including building, preparing your explanation, and recording the video. There is no extra credit for spending more time on your solution. The cap is tight on purpose. We want to see what you choose to leave out.
- Use your own tools, accounts, and AI tokens. Free-tier tools are completely fine. We are assessing judgment, not spend. KPMG will not provide tools, tokens, credentials, data, or environments for this exercise.
- Use only synthetic, public, or sample data. Do not use confidential or personal data.
- Make your artifact inspectable with the lightest supporting materials needed, such as a demo link, local files, screenshots, a README, sample outputs, or clear instructions.
- We expect you to use AI tools. This role does not exist without them. But remember: you can outsource thinking, you cannot outsource understanding. We will ask follow-up questions, and you will be expected to answer them. You should be able to disclose how you used AI and which decisions were yours.

What we will look for

- Understanding of the AI Builder role and the hiring challenge.
- Ability to turn ambiguity into a concrete direction, including how you frame the problem.



- A practical, usable artifact grounded in the job context.
- Evidence of builder mindset: making something real, not only describing an idea.
- Communication of assumptions, tradeoffs, risks, and next steps.
- Responsible AI judgment, including fairness, governance, transparency, and candidate experience.

We are not looking for polish for its own sake. We are looking for evidence that you can think in systems, build pragmatically, communicate clearly, and balance speed with judgment in a complex enterprise environment.

“Like many complex organizations, KPMG has competing priorities, a need for risk, security and appropriate governance and a desire to move faster and be innovative in the way we operate and deliver to clients - you are expected to consider the needs and sometime competing priorities of multiple stakeholders when considering how you approach answering this question”



APPENDIX A: AI Builder Job Description – Senior Consultant

Overview

At KPMG in Canada, our people bring their unique perspectives to Canada's most important challenges. Here, you can build momentum that reaches beyond our business, develop skills for the future, and take ownership of your career with support at every stage. Join a firm where your career can make a difference.

KPMG in Canada is scaling an internal team of **AI Builders** who operate as the engine embedded in the business—powering the firm's agent builds and transforming our business

This team is for people who sit at the intersection of enterprise operations, technology, and design, and who can look at how the firm works today, imagine how it should operate with agentic AI, and make that future real through execution.

Our objective is to move beyond experimentation and enable real internal adoption of agentic AI, redesigning core workflows, decision systems, and operating models using autonomous agents that power KPMG's service lines and internal teams.

Within this role, AI Builders take ownership of the full build lifecycle, working closely with Engineering and Internal Transformation teams to move solutions from our KPMG AI Lab into scalable, reliable production systems, all within a complex and regulated enterprise.

What you will do

At the Senior Consultant level on this team, you will be a hands-on AI Builder who takes ownership of well-scoped internal problems and contributes to building real, agentic solutions end to end. You will work closely with senior builders, practice leads, and internal stakeholders on rapid experimentation, prototyping, and Lab-based initiatives that demonstrate and validate new agentic capabilities.

In this role, you will:

- Build, test, and harden agents, plugins, skills, and tools using the firm's approved technology stack, working in deep connection with cloud and engineering teams to deliver production-grade solutions.
- Own the full build lifecycle: prototype in the AI Lab, productionize through Internal Transformation, and maintain solutions post-deployment with support from engineering and platform teams.
- Work with production systems, design evaluations, and handle integrations with enterprise platforms and data sources.
- Contribute reusable skills, components, and tools back into the shared AI stack so each subsequent build starts further ahead.
- Operate inside the shared builder pool, deployed against the highest-leverage agent builds across the organization.
- Be an active member of the AI Lab, contributing demos, reusable assets, and helping build a strong builder community across the firm.



What you bring to the role

At KPMG in Canada, we know the strongest builders come from a wide range of backgrounds, including engineering, business, strategy, design, and self-taught technical paths. This role is designed for early-career professionals, who are ready to build, learn, and grow quickly in a complex and regulated enterprise environment.

- You bring enterprise fluency and understand how organizations operate.
- You can identify real pain points, constraints, and opportunities where agentic AI creates tangible value.
- You think in workflows, not just models, designing multi-step processes that account for human-AI handoffs, decision boundaries, and appropriate levels of autonomy.
- You have a builder mindset and don't just describe solutions, you create them, whether through code, configuration, low-code platforms, APIs, or orchestration tools.
- You take end-to-end ownership, from problem definition through to a working system that people actually use.
- You're comfortable operating in ambiguity and understand that enterprise transformation is rarely linear.
- You can make reasonable assumptions, build quickly, test, learn, and iterate. You also apply sound judgment, treating risk, governance, ethics, and trust as core design constraints rather than afterthoughts.

Where and how you'll work

You will be based out of KPMG's AI Labs in Toronto, Vancouver, Calgary, or Montreal

As part of the internal builder model, your work will align primarily the following context:

Embedded in business units (AI Builder – Internal): Supporting internal teams by building and testing agentic workflows and prototypes that explore how autonomous agents could improve internal delivery and operations. You'll work alongside a core internal team and gain exposure to a range of internal use cases, tools, and applied research efforts across the firm.

What's in it for you

- The opportunity to work on genuinely new agentic AI use cases at the experimentation and prototyping stage.
- Exposure to complex, regulated, enterprise-scale internal problems and workflows.
- Hands-on experience building demos, prototypes, and showcase artifacts that influence how AI is adopted internally.
- A multidisciplinary environment with builders, researchers, designers, and business professionals working side by side.
- An AI Lab built for experimentation, with secure infrastructure, frontier models, and applied research.
- **Access to frontier models and platforms**, including early exposure to leading foundation models, agent frameworks, and enterprise-grade AI tooling through KPMG's secure AI Lab environment, enabling hands-on experimentation, prototyping, and applied learning at the forefront of agentic AI.