

AI Builder Candidate Assignment Submission

Name: Mansi Dhanania

Artifact: AI Builder Evidence Workbench

Demo link: <https://github.com/MansiDhanania/AI-Builder-Tool>

Video link: https://drive.google.com/file/d/1-SZ9fl5aKDVx65b_5V-ttgkPxdwqDxIU/view?usp=sharing

What I Built

I built a lightweight live-AI reviewer workbench for the AI Builder case-study stage. The tool helps a recruiter review open-ended candidate submissions such as resumes, code, portfolios, design write-ups, README files, and videos. Instead of creating a scoring system or automated hire/reject label, the system uses a real multi-agent LLM pipeline to call tools, map evidence to job-specific requirements and turn gaps into targeted interview questions.

Problem Framing

For this role, I prioritized extracting the evidence of:

- Ambiguity to Scope
- Agentic Workflow Thinking
- Builder Execution
- Responsible AI and Governance
- Enterprise Fluency
- Reusable Stack Contribution
- Communication and Tradeoffs, from the submitted documents.

Prototype Workflow

1. Recruiter selects a role.
2. Recruiter views candidates loaded from cleaned folders in candidate_submissions/.
3. A multi-agent LLM pipeline inspects available artifacts to produce transparent logs.
4. Resume is highlighted by category: technical stack, organizational knowledge and soft skills.
5. Evidence is mapped to submission document snippets with reasoning.
6. A constrained chat assistant answers clarification questions without giving hiring opinions.
7. The human reviewer records either "Recommend for Panel Interview" or "Flag for Additional Clarification."

Responsible AI Choices

- No numerical candidate score.
- No automated hire/reject recommendation.
- Synthetic sample data only.
- Missing optional artifacts are not treated as candidate failure.
- Assistant responses are constrained to submitted materials and the role rubric.
- Final accountability remains with the human reviewer.

AI Use

I used AI both during development and inside the artifact. The app backend calls Groq-hosted Llama models, runs a multi-agent pipeline and asks the model to synthesize source-grounded evidence. Azure was setup as the intended hosting platform but the project runs locally for demo purposes. The core design and implementation decisions were mine: evidence mapping over scoring, explicit human-AI boundaries, a constrained reviewer chat and interview-question generation from evidence gaps.

Tradeoffs and Next Steps

This prototype uses synthetic cleaned candidate folders so reviewers can inspect the workflow without confidential data. A production version would add raw PDF/DOCX/image parsing, video transcription, deeper repository scanning, automated redaction, evaluator calibration and enterprise deployment controls.